

```
In [1]: 1 import numpy as np
        2 import math
        3 import scipy.integrate
        4 import matplotlib.pyplot as plt
```

```
In [2]: 1 # Get floating-point information for default float data type
        2 float_info = np.finfo(float)
        3 # Machine epsilon
        4 print(float_info.eps)
```

2.220446049250313e-16

## Lab Book 01

```
In [3]: 1 # a)
        2 print("the largest positive finite machine number is", float_in
        3 # b)
        4 print("the largest absolute value of negative finite machine nu
        5 # c)
        6 print("smallest positive machine number is", float_info.tiny)
```

the largest positive finite machine number is 1.7976931348623157e+308

the largest absolute value of negative finite machine number is 1.7976931348623157e+308

smallest positive machine number is 2.2250738585072014e-308

```
In [4]: 1 # Generate np.nan
2 x = np.arange(5)
3 print(x / x)
4 # Generate np.inf
5 y = np.ones((5,))
6 print(y / x)
7 # Is infinity equal to itself?
8 print(np.inf == np.inf)
9 # Is nan equal to itself?
10 print(np.nan == np.nan)
11 # Compare nan to other values
12 print(np.nan < 0, np.nan == 0, np.nan > 0)
13
```

```
[nan  1.  1.  1.  1.]
[      inf  1.          0.5          0.33333333  0.25          ]
True
False
False False False
```

/Users/x\_x/opt/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:3: RuntimeWarning: invalid value encountered in true\_divide

This is separate from the ipykernel package so we can avoid doing imports until

/Users/x\_x/opt/anaconda3/lib/python3.7/site-packages/ipykernel\_launcher.py:6: RuntimeWarning: divide by zero encountered in true\_divide

```
In [5]: 1 print(1+2==3)
2 print(0.1+0.2==0.3)
3 print(0.1+0.2+0.7==1)
```

```
True
False
True
```

## Lab Book 02

```
In [6]: 1 print("sin(pi) = ",math.sin(math.pi))
2 print("cos(pi/2) = ",math.cos(0.5*math.pi))
3 print("sin(pi/2) = ",math.sin(0.5*math.pi))
4 print("cos(0) = ",math.cos(0))
5 print("cos(pi) = ",math.cos(math.pi))
```

```
sin(pi) =  1.2246467991473532e-16
cos(pi/2) =  6.123233995736766e-17
sin(pi/2) =  1.0
cos(0) =  1.0
cos(pi) = -1.0
```

When the value of sin or cos should be 0, python can't return the exact value but can only return extremely small number. When the value should be +-1, python can give the exactly right answer.

## Lab Book 03

```
In [7]: 1 def f(x):  
2         return (np.cos(2*x)-1) / pow(x,2)  
3 def taylorF(x):  
4         return 2 * pow(x,2)/3 - 2  
5 def alterF(x):  
6         return -2 * pow(np.sin(x),2) / pow(x,2)
```

```
In [8]: 1 x=np.array([1e-2,1e-3,1e-4,1e-5,1e-6,1e-7,1e-8,1e-9,1e-10])
2 for item in x:
3     print("x = ", item)
4     print("f(x) = ", f(item))
5     print("result of Tolor series = ", taylorF(item), ", the er
6     print("improved result = ", alterF(item), "the difference w
```

```
x = 0.01
f(x) = -1.9999333342224368
result of Tolor series = -1.9999333333333333 , the error is 8.89
1034575242429e-10
improved result = -1.9999333342222159 the difference with Taylor
series is 8.888825231423425e-10
x = 0.001
f(x) = -1.999999333368585
result of Tolor series = -1.9999993333333332 , the error is 3.52
5180147789797e-11
improved result = -1.9999993333334223 the difference with Taylor
series is 8.903988657493755e-14
x = 0.0001
f(x) = -1.999999987845058
result of Tolor series = -1.9999999933333332 , the error is 5.48
8275167664369e-09
improved result = -1.9999999933333334 the difference with Taylor
series is 2.220446049250313e-16
x = 1e-05
f(x) = -2.0000001654807416
result of Tolor series = -1.9999999999333333 , the error is 1.65
54740822627423e-07
improved result = -1.9999999999333333 the difference with Taylor s
eries is 2.220446049250313e-16
x = 1e-06
f(x) = -1.999955756559757
result of Tolor series = -1.9999999999993334 , the error is 4.42
4343957643018e-05
improved result = -1.9999999999993332 the difference with Taylor
series is 2.220446049250313e-16
x = 1e-07
f(x) = -1.998401444325282
result of Tolor series = -1.9999999999999933 , the error is 0.00
15985556747113439
improved result = -1.9999999999999933 the difference with Taylor
series is 0.0
x = 1e-08
f(x) = -2.2204460492503126
result of Tolor series = -2.0 , the error is 0.22044604925031264
improved result = -2.0 the difference with Taylor series is 0.0
x = 1e-09
f(x) = 0.0
result of Tolor series = -2.0 , the error is 2.0
improved result = -2.0 the difference with Taylor series is 0.0
x = 1e-10
f(x) = 0.0
result of Tolor series = -2.0 , the error is 2.0
improved result = -2.0 the difference with Taylor series is 0.0
```

Python has a bigger error when  $x$  is smaller when calculating  $f(x)$  because of the subtraction in the numerator will lead to catastrophic cancellation when two values are close to 0 and therefore leads to big error.

## **Lab Book 04**

In [9]:

```
1 def expTaylor(x):
2     term = 1.0
3     result = 0.0
4     n = 1
5     while term + result != result:
6         term = pow(x,n) / math.factorial(n)
7         result += term
8         n += 1
9     return result + 1
10 for item in x:
11     print("x = ", item)
12     print("exp(x) = ",math.exp(item)-1)
13     print("expm1(x) = ",math.expm1(item))
14     print("Taylor series = ",expTaylor(item)-1)
```

```
x = 0.01
exp(x) = 0.010050167084167949
expm1(x) = 0.010050167084168058
Taylor series = 0.010050167084167949
x = 0.001
exp(x) = 0.0010005001667083846
expm1(x) = 0.0010005001667083417
Taylor series = 0.0010005001667083846
x = 0.0001
exp(x) = 0.0001000050001667141
expm1(x) = 0.00010000500016667084
Taylor series = 0.0001000050001667141
x = 1e-05
exp(x) = 1.0000050000069649e-05
expm1(x) = 1.0000050000166668e-05
Taylor series = 1.0000050000069649e-05
x = 1e-06
exp(x) = 1.0000004999621837e-06
expm1(x) = 1.0000005000001665e-06
Taylor series = 1.0000004999621837e-06
x = 1e-07
exp(x) = 1.0000000494336803e-07
expm1(x) = 1.0000000500000016e-07
Taylor series = 1.0000000494336803e-07
x = 1e-08
exp(x) = 9.9999993922529e-09
expm1(x) = 1.0000000050000001e-08
Taylor series = 9.9999993922529e-09
x = 1e-09
exp(x) = 1.000000082740371e-09
expm1(x) = 1.0000000005000001e-09
Taylor series = 1.000000082740371e-09
x = 1e-10
exp(x) = 1.000000082740371e-10
expm1(x) = 1.000000000005e-10
Taylor series = 1.000000082740371e-10
```

Since  $e^x$  is quite close to 1 when  $x$  is close to 0, therefore the digits after 1 can not be stored properly and when calculating  $e^x - 1$ , the catastrophic cancellation will occur and lead to inaccurate results compared to `expm1(x)`. Similar error occurs when using Taylor series expansion due to the same result.

## Lab Book 05

In [10]:

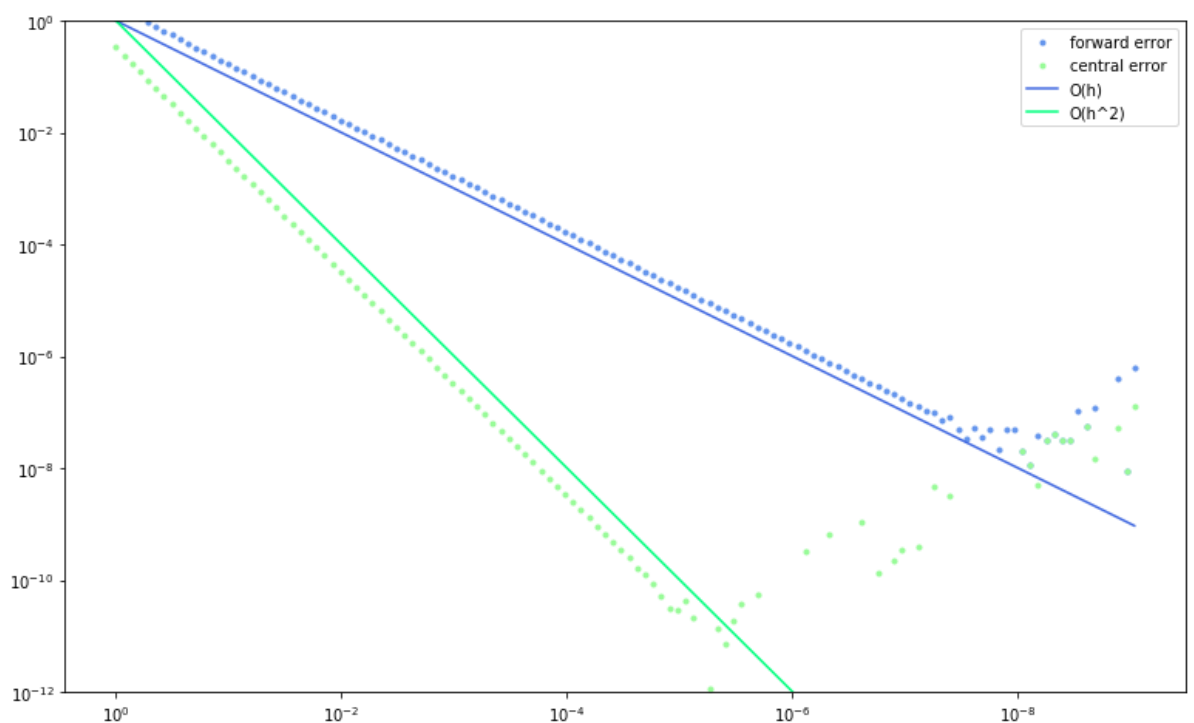
```
1 def f(x):  
2     return math.exp(x) - math.cos(x)  
3 def df(x):  
4     return math.exp(x) + math.sin(x)
```

```

In [11]: 1 # Decreasing sequence of h values
2 hs = 2**(-np.linspace(0,30,128))
3 # Make an empty vector of errors
4 forward_error = np.zeros((len(hs),))
5 central_error = np.zeros((len(hs),))
6 for i in range(len(hs)):
7     h = hs[i] # current value of h
8     forward_error[i] = (f(1+h) - f(1)) / h - df(1)
9     central_error[i] = (f(1+h) - f(1-h)) / (2*h) - df(1)
10 # The below vectors show the error decrease for particular orde
11 first_order_rate = hs # plot hs vs first_order_rate to see
12 second_order_rate = hs**2 # plot hs vs second_order_rate to se
13
14 # Plot hs vs forward_error and central_error (on a log-log plot
15 plt.figure(figsize = (13,8))
16 plt.clf()
17 plt.loglog(hs, forward_error, '.', label = "forward error", colo
18 plt.loglog(hs, central_error, '.', label = "central error", colo
19 plt.loglog(hs, first_order_rate, label = "O(h)", color = 'royal
20 plt.loglog(hs, second_order_rate, label = "O(h^2)", color='spri
21 plt.legend(loc = 'best')
22 ax = plt.gca()
23 ax.invert_xaxis()
24 plt.ylim(1e-12,1)
25 plt.show

```

Out[11]: <function matplotlib.pyplot.show(\*args, \*\*kw)>



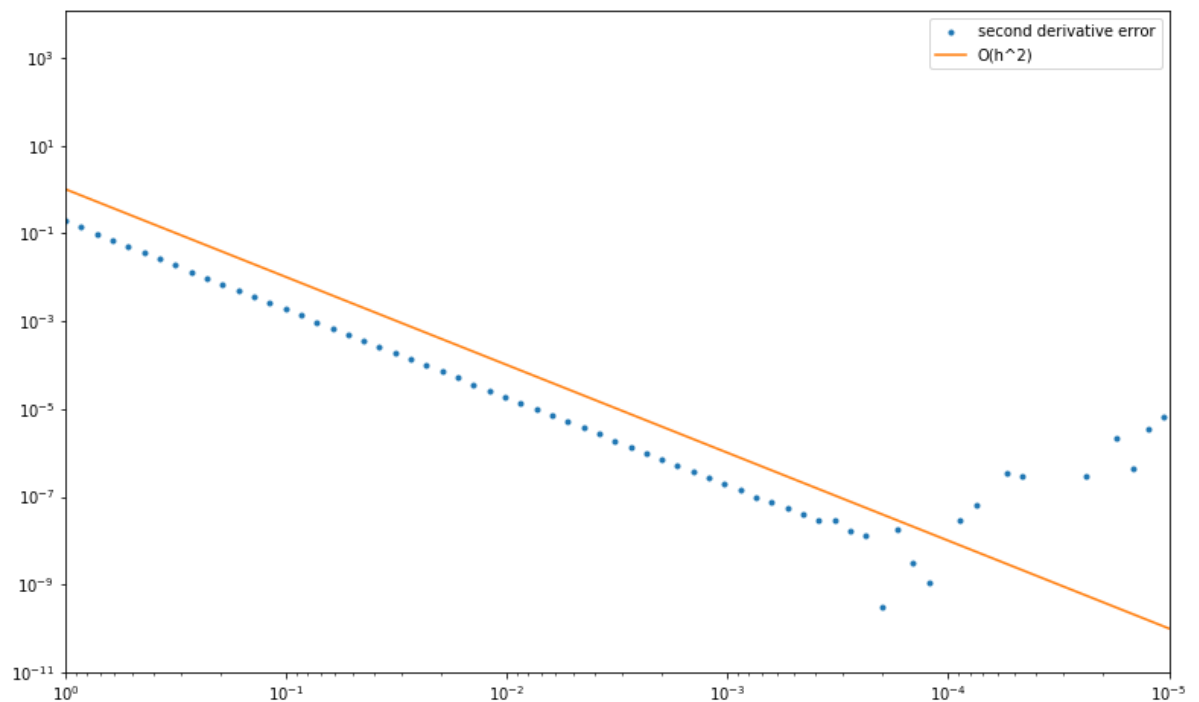
As  $h$  approaches 0, the forward error is almost the same as  $O(h)$ 's convergence rate before reaching the machine epsilon, which suggests it's order of convergence is 1. Similarly, the central error is almost the same as  $O(h^2)$  before reaching the machine epsilon, therefore it has an order of convergence of 2.



## Lab Book 06

```
In [12]: 1 def d2f(x):
2         return math.exp(x) + math.cos(x)
3 second_error = np.zeros((len(hs),))
4 for i in range(len(hs)):
5     h = hs[i]
6     second_error[i] = (f(1+h) - 2*f(1) + f(1-h)) / pow(h,2) - d
7
8 plt.figure(figsize = (13,8))
9 plt.clf()
10 plt.loglog(hs, second_error, '.', label = "second derivative error")
11 plt.loglog(hs, second_order_rate, label = "O(h^2)")
12 plt.legend(loc = 'best')
13 plt.xlim(1e-5,1)
14 plt.ylim(1e-11,)
15 ax = plt.gca()
16 ax.invert_xaxis()
17 plt.show
```

```
Out[12]: <function matplotlib.pyplot.show(*args, **kw)>
```



As  $h$  approaches to 0, the error of second order derivative's approximation is almost the same as  $O(h^2)$ 's convergence rate before reaching the machine epsilon, this suggests that it's order of convergence is 2.

## Lab Book 07

```
In [13]: 1 def left_riemann_sum(xs, ys):
2         n = len(xs)
3         if len(xs) != len(ys):
4             return np.NaN
5         integrals = np.zeros(n)
6         x0s = xs[:n-1]
7         x1s = xs[1:]
8         integrals = (x1s-x0s) * ys[:n-1]
9         return sum(integrals)
10 def f(x):
11     return math.exp(-x)
```

```
In [14]: 1 a = -1
2         b = 1
3         n = 100
4         xs = np.linspace(a, b, n) # equally spaced points xi
5         ys = np.zeros(n)
6         for i in range(n-1):
7             ys[i] = f(xs[i])
8         print("Left Riemann sum =", left_riemann_sum(xs, ys))
9         print("True value is around 2.350402387287603")
```

```
Left Riemann sum = 2.374223762501847
True value is around 2.350402387287603
```

The left Reimann Sum can roughly approximate the integral and has 2 acutare digits. The accuracy can also be improved by adding more points and make each points closer but in a raletively low convergence rate.

## Lab Book 08

```
In [15]: 1 def trapezoidal_rule(xs, ys):
2         n = len(xs)
3         if len(xs) != len(ys):
4             return np.NaN
5         integrals = np.zeros(n)
6         x0s = xs[:n-1]
7         x1s = xs[1:]
8         y0s = ys[:n-1]
9         y1s =ys[1:]
10        integrals = (x1s-x0s) * (y0s+y1s) / 2
11        return sum(integrals)
```

```
In [16]: 1 print("Trapezoidal Rule =", trapezoidal_rule(xs, ys))
2         print("Left Riemann sum =", left_riemann_sum(xs, ys))
3         print("True value is around 2.350402387287603")
```

```
Trapezoidal Rule = 2.346766370295189
Left Riemann sum = 2.374223762501847
True value is around 2.350402387287603
```

The Trapezoidal Rule clearly has a much more accuracy with same number of points although it also can only give two accurate digits. When increasing the number of points, the Trapezoidal Rule has a faster convergence rate than left Riemann Sum.

## Lab Book 09

```
In [17]: 1 def runges(x):  
2         return 1 / (1 + 25*pow(x,2))
```

```
In [18]: 1 n = 20  
2 xs = np.linspace(a, b, n) # equally spaced points xi  
3 ys = runges(xs)  
4 print("Simpson's Rule = ", scipy.integrate.simps(ys, xs))  
5 print("Gaussian Quadrature = ", scipy.integrate.fixed_quad(rung  
6 print("True value is around 0.549360307")
```

```
Simpson's Rule = 0.5493758748195977  
Gaussian Quadrature = 0.548997098104952  
True value is around 0.549360307
```

Both ways gives three accurate digits compared to the true value. The Simpson's Rule in python is more accurate than Gaussian Quadrature and have errors at  $10^{-5}$  and  $10^{-3}$  respectively.

## Lab Book 10

```
In [19]: 1 print("The approximation is ", scipy.integrate.quad(runges, a,
```

```
The approximation is 0.5493603067780066 , the error is 2.8668279  
350011863e-09
```

The output is quite accurate compared to the true value and the error is at  $10^{-9}$ , which is a very small error compared to the above two methods.

```
In [ ]: 1
```